


# CS 188: Artificial Intelligence


## Perceptrons and Logistic Regression



Pieter Abbeel & Dan Klein  
University of California, Berkeley


## Linear Classifiers

---



# Feature Vectors


---



## Some (Simplified) Biology

---


- Very loose inspiration: human neurons



# Linear Classifiers


---

- Inputs are **feature values**
- Each feature has a **weight**
- Sum is the **activation**



$$\text{activation}_w(x) = \sum_i w_i \cdot f_i(x) = w \cdot f(x)$$


- If the activation is:
  - Positive, output +1
  - Negative, output -1



## Weights


---

- Binary case: compare features to a weight vector
- Learning: figure out the weight vector from examples



# Decision Rules

---





## Binary Decision Rule

---

- In the space of feature vectors
  - Examples are points
  - Any weight vector is a hyperplane
  - One side corresponds to  $Y=+1$
  - Other corresponds to  $Y=-1$

$w$

BIAS	:	-3
free	:	4
money	:	2
...		



# Weight Updates


---




## Learning: Binary Perceptron

---


- Start with weights = 0
- For each training instance:
  - Classify with current weights



- If correct (i.e.,  $y=y^*$ ), no change!



- If wrong: adjust the weight vector



# Learning: Binary Perceptron


---

- Start with weights = 0
- For each training instance:
  - Classify with current weights

$$y = \begin{cases} +1 & \text{if } w \cdot f(x) \geq 0 \\ -1 & \text{if } w \cdot f(x) < 0 \end{cases}$$

- If correct (i.e.,  $y=y^*$ ), no change!
- If wrong: adjust the weight vector by adding or subtracting the feature vector. Subtract if  $y^*$  is -1.


$$w = w + y^* \cdot f$$



## Examples: Perceptron

---

- Separable Case



# Multiclass Decision Rule

---

- If we have multiple classes:

- A weight vector for each class:


$$w_y$$

- Score (activation) of a class  $y$ :


$$w_y \cdot f(x)$$

- Prediction highest score wins

$$y = \arg \max_y w_y \cdot f(x)$$



$w_1 \cdot f$  biggest



Binary = multiclass where the negative class has weight zero

## Learning: Multiclass Perceptron

---


- Start with all weights = 0
- Pick up training examples one by one
- Predict with current weights

$$y = \arg \max_y w_y \cdot f(x)$$

- If correct, no change!
- If wrong: lower score of wrong answer, raise score of right answer

$$w_y = w_y - f(x)$$

$$w_{y^*} = w_{y^*} + f(x)$$



## Example: Multiclass Perceptron

---

“win the vote”

“win the election”

“win the game”

$w_{SPORTS}$

BIAS	:	1
win	:	0
game	:	0
vote	:	0
the	:	0
...		

$w_{POLITICS}$

BIAS	:	0
win	:	0
game	:	0
vote	:	0
the	:	0
...		

$w_{TECH}$

BIAS	:	0
win	:	0
game	:	0
vote	:	0
the	:	0
...		


## Properties of Perceptrons

---


- Separability: true if some parameters get the training set perfectly correct
- Convergence: if the training is separable, perceptron will eventually converge (binary case)
- Mistake Bound: the maximum number of mistakes (binary case) related to the *margin* or degree of separability

$$\text{mistakes} < \frac{k}{\delta^2}$$

Separable





Non-Separable




# Problems with the Perceptron

---



- Noise: if the data isn't separable, weights might thrash
  - Averaging weight vectors over time can help (averaged perceptron)



- Mediocre generalization: finds a "barely" separating solution




- Overtraining: test / held-out accuracy usually rises, then falls
  - Overtraining is a kind of overfitting




# Improving the Perceptron

---




## Non-Separable Case: Deterministic Decision

---



## Non-Separable Case: Probabilistic Decision

---




# How to get probabilistic decisions?

---

- Perceptron scoring:  $z = w \cdot f(x)$
- If  $z = w \cdot f(x)$  very positive → want probability going to 1
- If  $z = w \cdot f(x)$  very negative → want probability going to 0
- Sigmoid function

$$\phi(z) = \frac{1}{1 + e^{-z}}$$



## Best w?

---

- Maximum likelihood estimation:

$$\max_w ll(w) = \max_w \sum_i \log P(y^{(i)} | x^{(i)}; w)$$



with:

$$P(y^{(i)} = +1 | x^{(i)}; w) = \frac{1}{1 + e^{-w \cdot f(x^{(i)})}}$$
$$P(y^{(i)} = -1 | x^{(i)}; w) = 1 - \frac{1}{1 + e^{-w \cdot f(x^{(i)})}}$$

= Logistic Regression



## Separable Case: Deterministic Decision – Many Options

---



## Separable Case: Probabilistic Decision – Clear Preference

---




# Multiclass Logistic Regression

---

- Recall Perceptron:

- A weight vector for each class:  $w_y$
- Score (activation) of a class  $y$ :  $w_y \cdot f(x)$
- Prediction highest score wins  $y = \arg \max_y w_y \cdot f(x)$



- How to make the scores into probabilities?

$$z_1, z_2, z_3 \rightarrow \underbrace{\frac{e^{z_1}}{e^{z_1} + e^{z_2} + e^{z_3}}, \frac{e^{z_2}}{e^{z_1} + e^{z_2} + e^{z_3}}, \frac{e^{z_3}}{e^{z_1} + e^{z_2} + e^{z_3}}}_{\text{softmax activations}}$$

original activations

## Best $w$ ?

---

- Maximum likelihood estimation:

$$\max_w ll(w) = \max_w \sum_i \log P(y^{(i)} | x^{(i)}; w)$$

with:  $P(y^{(i)} | x^{(i)}; w) = \frac{e^{w_{y^{(i)}} \cdot f(x^{(i)})}}{\sum_y e^{w_y \cdot f(x^{(i)})}}$

= Multi-Class Logistic Regression

## Next Lecture

---

- Optimization

- i.e., how do we solve:

$$\max_w \text{ll}(w) = \max_w \sum_i \log P(y^{(i)}|x^{(i)}; w)$$